

Entrene un autoencoder para obtener una representación de baja dimensionalidad de las imágenes de MNIST. Use dichas representaciones para entrenar un perceptrón multicapa como clasificador. ¿Cuál es el tiempo de entrenamiento y la exactitud del clasificador obtenido cuando parte de la representación del autoencoder, en comparación con lo obtenido usando las imágenes originales?

En este caso se nos permitió usar Pytorch para la parte de código. La idea es definir la estructura general de la red y luego ir modificandola hasta lograr la mejor performance.

```
In [7]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
import torch.optim as optim
import matplotlib.pyplot as plt

# Elegir dispositivo
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)

# Transformación
transform = transforms.ToTensor()

# Dataset MNIST
train_dataset = datasets.MNIST(root="data", train=True, download=True, transform=transform)
test_dataset = datasets.MNIST(root="data", train=False, download=True, transform=transform)

# minibatch
batch_size = 64
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=len(test_dataset))
```

Device: cuda

Primero una función que entrena y grafica

```
In [8]: def train_and_evaluate(model, train_loader, test_loader,
                           criterion, optimizer, device,
                           num_epochs=50, patience=5, min_delta=1e-4):

    model = model.to(device)
    best_loss = float("inf")
    epochs_no_improve = 0

    history = {"train_loss": [], "test_loss": [], "accuracy": []}

    total_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
    print(model)
    print(f"Parámetros entrenables: {total_params:,}")

    for epoch in range(num_epochs):
        model.train()
        train_loss = 0.0

        for x,y in train_loader:
            x, y = x.to(device), y.to(device)
            optimizer.zero_grad()
            outputs = model(x)
            loss = criterion(outputs, y)
            loss.backward()
```

```
optimizer.step()
train_loss += loss.item()

train_loss /= len(train_loader)
history["train_loss"].append(train_loss)

# Evaluación
model.eval()
test_loss = 0.0
correct, total = 0, 0

with torch.no_grad():
    for x,y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = model(x)
        test_loss += criterion(outputs, y).item()
        _, pred = outputs.max(1)
        total += y.size(0)
        correct += pred.eq(y).sum().item()

test_loss /= len(test_loader)
accuracy = 100 * correct / total

history["test_loss"].append(test_loss)
history["accuracy"].append(accuracy)

print(f"Epoch {epoch+1} | Train Loss {train_loss:.4f} | "
      f"Test Loss {test_loss:.4f} | Acc {accuracy:.2f}%")

# Early stopping
if test_loss < best_loss - min_delta:
    best_loss = test_loss
    epochs_no_improve = 0
else:
    epochs_no_improve += 1

if epochs_no_improve >= patience:
    print("✓ Early stopping activado")
    break

# Gráficos
plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
plt.plot(history["train_loss"], label="Train")
plt.plot(history["test_loss"], label="Test")
plt.title("Loss")
plt.legend()

plt.subplot(1,2,2)
plt.plot(history["accuracy"], label="Accuracy")
plt.title("Accuracy (%)")
plt.legend()
plt.show()

return history
```

In [9]:

```
def train_autoencoder(model, train_loader, test_loader, criterion, optimizer, device, num_epochs=10):
    train_losses = []
    test_losses = []
    model.to(device)
```

```

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    for x, _ in train_loader:
        x = x.to(device)
        x_flat = x.view(x.size(0), -1)      # aplana
        optimizer.zero_grad()
        outputs = model(x_flat)
        loss = criterion(outputs, x_flat)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    train_losses.append(running_loss / len(train_loader))

    # evaluación
    model.eval()
    test_loss = 0.0
    with torch.no_grad():
        for x, _ in test_loader:
            x = x.to(device)
            x_flat = x.view(x.size(0), -1)
            outputs = model(x_flat)
            loss = criterion(outputs, x_flat)
            test_loss += loss.item()
    test_losses.append(test_loss / len(test_loader))

    print(f"Epoch [{epoch+1}/{num_epochs}] - Train Loss: {train_losses[-1]:.4f}, Test Loss: {test_losses[-1]:.4f}")

return train_losses, test_losses

```

La primer estructura que quiero hacer es el autoencoder. Vo a elegir una arquitectura simétrica con una reducción del 80% de tamaño de dimensión de entrada, de acuerdo al principio de Pareto, que dice que el 80% de las cosas son causadas por el 20%. Esto lo extrapoló a energía y varianza y listo.

Entramos con 28×28 , luego a $28 \times 28 \times 0.5$ luego a $0.2 \times 28 \times 28$. De ahí sube de vuelta

```
In [10]: class Autoencoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(28*28, 392)
        self.fc2 = nn.Linear(392, 156)
        self.fc3 = nn.Linear(156, 392)
        self.fc4 = nn.Linear(392, 28*28)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = torch.sigmoid(self.fc4(x))
        return x

model_cnn = Autoencoder().to(device)
```

```
In [11]: # ----- Entrenamiento -----
criterion = nn.BCELoss()
optimizer = optim.Adam(model_cnn.parameters(), lr=1e-3)
```

```
train_losses, test_losses = train_autoencoder(
    model_cnn, train_loader, test_loader,
    criterion, optimizer, device,
    num_epochs=100
)

# ----- Gráficos -----
plt.figure(figsize=(8,4))
plt.plot(train_losses, label='Train Loss')
plt.plot(test_losses, label='Test Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('Pérdida del Autoencoder')
plt.show()

# ----- Visualizar reconstrucciones -----
model_cnn.eval()
with torch.no_grad():
    for x, _ in test_loader:
        x = x.to(device)
        x_flat = x.view(x.size(0), -1)
        outputs = model_cnn(x_flat)
        break # solo un batch

# convertir a imágenes
x = x.cpu().view(-1, 1, 28, 28)
outputs = outputs.cpu().view(-1, 1, 28, 28)

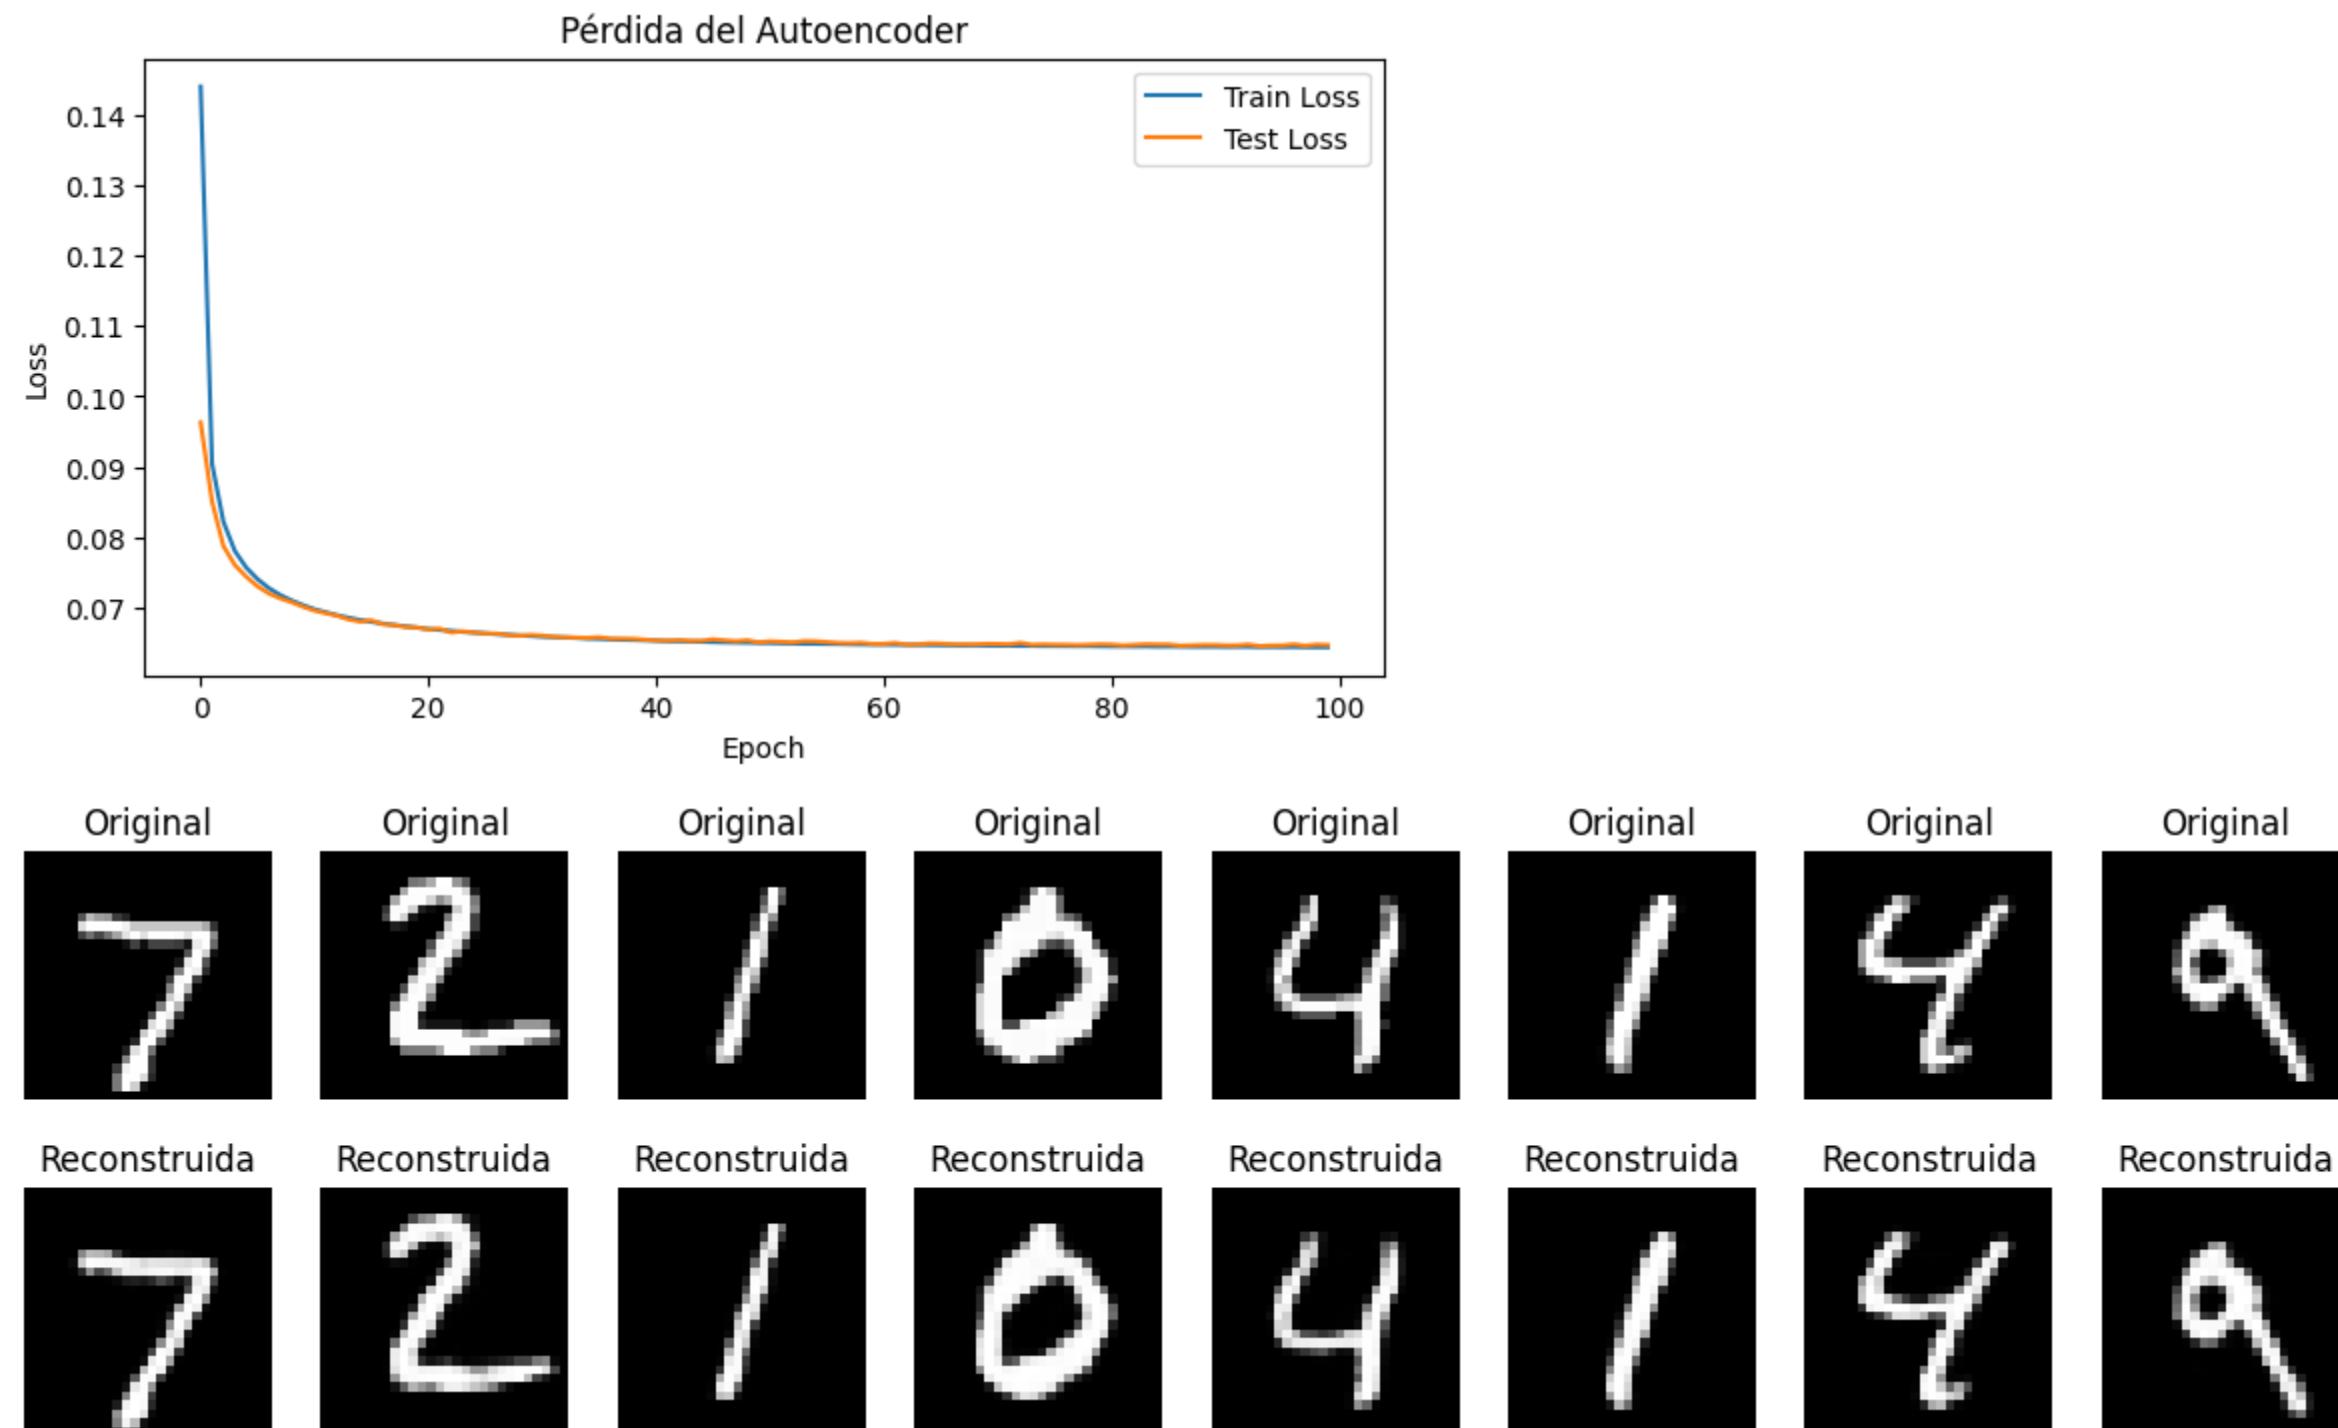
# mostrar algunas imágenes
n = 8
plt.figure(figsize=(15,4))
for i in range(n):
    # original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x[i][0], cmap="gray")
    plt.title("Original")
    plt.axis("off")

    # reconstruida
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(outputs[i][0], cmap="gray")
    plt.title("Reconstruida")
    plt.axis("off")

plt.show()
```

Epoch [1/100] - Train Loss: 0.1440, Test Loss: 0.0963
Epoch [2/100] - Train Loss: 0.0904, Test Loss: 0.0851
Epoch [3/100] - Train Loss: 0.0822, Test Loss: 0.0788
Epoch [4/100] - Train Loss: 0.0781, Test Loss: 0.0761
Epoch [5/100] - Train Loss: 0.0757, Test Loss: 0.0744
Epoch [6/100] - Train Loss: 0.0741, Test Loss: 0.0730
Epoch [7/100] - Train Loss: 0.0728, Test Loss: 0.0720
Epoch [8/100] - Train Loss: 0.0718, Test Loss: 0.0713
Epoch [9/100] - Train Loss: 0.0710, Test Loss: 0.0708
Epoch [10/100] - Train Loss: 0.0704, Test Loss: 0.0701
Epoch [11/100] - Train Loss: 0.0698, Test Loss: 0.0696
Epoch [12/100] - Train Loss: 0.0694, Test Loss: 0.0692
Epoch [13/100] - Train Loss: 0.0689, Test Loss: 0.0689
Epoch [14/100] - Train Loss: 0.0686, Test Loss: 0.0683
Epoch [15/100] - Train Loss: 0.0683, Test Loss: 0.0680
Epoch [16/100] - Train Loss: 0.0680, Test Loss: 0.0682
Epoch [17/100] - Train Loss: 0.0678, Test Loss: 0.0677
Epoch [18/100] - Train Loss: 0.0675, Test Loss: 0.0675
Epoch [19/100] - Train Loss: 0.0674, Test Loss: 0.0673
Epoch [20/100] - Train Loss: 0.0672, Test Loss: 0.0672
Epoch [21/100] - Train Loss: 0.0670, Test Loss: 0.0669
Epoch [22/100] - Train Loss: 0.0669, Test Loss: 0.0670
Epoch [23/100] - Train Loss: 0.0667, Test Loss: 0.0665
Epoch [24/100] - Train Loss: 0.0666, Test Loss: 0.0667
Epoch [25/100] - Train Loss: 0.0665, Test Loss: 0.0665
Epoch [26/100] - Train Loss: 0.0664, Test Loss: 0.0663
Epoch [27/100] - Train Loss: 0.0663, Test Loss: 0.0663
Epoch [28/100] - Train Loss: 0.0662, Test Loss: 0.0661
Epoch [29/100] - Train Loss: 0.0661, Test Loss: 0.0660
Epoch [30/100] - Train Loss: 0.0660, Test Loss: 0.0662
Epoch [31/100] - Train Loss: 0.0659, Test Loss: 0.0661
Epoch [32/100] - Train Loss: 0.0658, Test Loss: 0.0659
Epoch [33/100] - Train Loss: 0.0658, Test Loss: 0.0658
Epoch [34/100] - Train Loss: 0.0657, Test Loss: 0.0658
Epoch [35/100] - Train Loss: 0.0656, Test Loss: 0.0657
Epoch [36/100] - Train Loss: 0.0656, Test Loss: 0.0658
Epoch [37/100] - Train Loss: 0.0655, Test Loss: 0.0656
Epoch [38/100] - Train Loss: 0.0655, Test Loss: 0.0656
Epoch [39/100] - Train Loss: 0.0654, Test Loss: 0.0656
Epoch [40/100] - Train Loss: 0.0654, Test Loss: 0.0655
Epoch [41/100] - Train Loss: 0.0653, Test Loss: 0.0654
Epoch [42/100] - Train Loss: 0.0653, Test Loss: 0.0653
Epoch [43/100] - Train Loss: 0.0652, Test Loss: 0.0654
Epoch [44/100] - Train Loss: 0.0652, Test Loss: 0.0653
Epoch [45/100] - Train Loss: 0.0652, Test Loss: 0.0653
Epoch [46/100] - Train Loss: 0.0651, Test Loss: 0.0655
Epoch [47/100] - Train Loss: 0.0651, Test Loss: 0.0654
Epoch [48/100] - Train Loss: 0.0651, Test Loss: 0.0653
Epoch [49/100] - Train Loss: 0.0651, Test Loss: 0.0654
Epoch [50/100] - Train Loss: 0.0650, Test Loss: 0.0651
Epoch [51/100] - Train Loss: 0.0650, Test Loss: 0.0652
Epoch [52/100] - Train Loss: 0.0650, Test Loss: 0.0652
Epoch [53/100] - Train Loss: 0.0649, Test Loss: 0.0651
Epoch [54/100] - Train Loss: 0.0649, Test Loss: 0.0653
Epoch [55/100] - Train Loss: 0.0649, Test Loss: 0.0652
Epoch [56/100] - Train Loss: 0.0649, Test Loss: 0.0651
Epoch [57/100] - Train Loss: 0.0649, Test Loss: 0.0650
Epoch [58/100] - Train Loss: 0.0648, Test Loss: 0.0650
Epoch [59/100] - Train Loss: 0.0648, Test Loss: 0.0650
Epoch [60/100] - Train Loss: 0.0648, Test Loss: 0.0649

```
Epoch [61/100] - Train Loss: 0.0648, Test Loss: 0.0649
Epoch [62/100] - Train Loss: 0.0648, Test Loss: 0.0650
Epoch [63/100] - Train Loss: 0.0647, Test Loss: 0.0648
Epoch [64/100] - Train Loss: 0.0647, Test Loss: 0.0648
Epoch [65/100] - Train Loss: 0.0647, Test Loss: 0.0650
Epoch [66/100] - Train Loss: 0.0647, Test Loss: 0.0649
Epoch [67/100] - Train Loss: 0.0647, Test Loss: 0.0648
Epoch [68/100] - Train Loss: 0.0647, Test Loss: 0.0648
Epoch [69/100] - Train Loss: 0.0647, Test Loss: 0.0648
Epoch [70/100] - Train Loss: 0.0646, Test Loss: 0.0649
Epoch [71/100] - Train Loss: 0.0646, Test Loss: 0.0649
Epoch [72/100] - Train Loss: 0.0646, Test Loss: 0.0648
Epoch [73/100] - Train Loss: 0.0646, Test Loss: 0.0650
Epoch [74/100] - Train Loss: 0.0646, Test Loss: 0.0647
Epoch [75/100] - Train Loss: 0.0646, Test Loss: 0.0648
Epoch [76/100] - Train Loss: 0.0646, Test Loss: 0.0648
Epoch [77/100] - Train Loss: 0.0646, Test Loss: 0.0648
Epoch [78/100] - Train Loss: 0.0646, Test Loss: 0.0647
Epoch [79/100] - Train Loss: 0.0646, Test Loss: 0.0648
Epoch [80/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [81/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [82/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [83/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [84/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [85/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [86/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [87/100] - Train Loss: 0.0645, Test Loss: 0.0646
Epoch [88/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [89/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [90/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [91/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [92/100] - Train Loss: 0.0645, Test Loss: 0.0647
Epoch [93/100] - Train Loss: 0.0645, Test Loss: 0.0648
Epoch [94/100] - Train Loss: 0.0644, Test Loss: 0.0646
Epoch [95/100] - Train Loss: 0.0644, Test Loss: 0.0646
Epoch [96/100] - Train Loss: 0.0644, Test Loss: 0.0647
Epoch [97/100] - Train Loss: 0.0644, Test Loss: 0.0648
Epoch [98/100] - Train Loss: 0.0644, Test Loss: 0.0646
Epoch [99/100] - Train Loss: 0.0644, Test Loss: 0.0648
Epoch [100/100] - Train Loss: 0.0644, Test Loss: 0.0647
```



1570 parámetros no congelados, la capa de 156 a 10 + bias

```
In [ ]: class ClassifierFromEncoder(nn.Module):
    def __init__(self, autoencoder, num_classes=10):
        super().__init__()
        # acá tomamos las capas del otro modelo y las congelamos
        self.encoder_fc1 = autoencoder.fc1
        self.encoder_fc2 = autoencoder.fc2
        for param in self.encoder_fc1.parameters():
            param.requires_grad = False
        for param in self.encoder_fc2.parameters():
            param.requires_grad = False
        self.classifier = nn.Linear(156, num_classes) # Lo más simple posible

    def forward(self, x):
        x = x.view(x.size(0), -1)
```

```
with torch.no_grad():
    x = F.relu(self.encoder_fc1(x))
    encoded = F.relu(self.encoder_fc2(x))
    logits = self.classifier(encoded)
    return logits

classifier_model = ClassifierFromEncoder(model_cnn).to(device)
criterion_cls = nn.CrossEntropyLoss()
optimizer_cls = optim.Adam(classifier_model.classifier.parameters(), lr=1e-3)

# -----
# Entrenamiento Clasificador
# -----
num_epochs_cls = 100
for epoch in range(num_epochs_cls):
    classifier_model.train()
    running_loss = 0
    correct, total = 0, 0
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)
        optimizer_cls.zero_grad()
        outputs = classifier_model(x)
        loss = criterion_cls(outputs, y)
        loss.backward()
        optimizer_cls.step()
        running_loss += loss.item()
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
    train_acc = 100 * correct / total
    print(f"Epoch {epoch+1}/{num_epochs_cls} | Loss={running_loss/len(train_loader):.4f} | Acc={train_acc:.2f}%")

# -----
# Evaluación en Test
# -----
classifier_model.eval()
correct, total = 0, 0
with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
print(f"Test Accuracy: {100 * correct/total:.2f}%")

# Evaluación por dígito
classifier_model.eval()
confusion_matrix = torch.zeros(10, 10)
digit_totals = torch.zeros(10)
digit_correct = torch.zeros(10)
total_correct = 0
total_samples = 0

with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
```

```
_> pred = outputs.max(1)

# Contadores globales
total_correct += (pred == y).sum().item()
total_samples += y.size(0)

# Contadores por dígito
for t, p in zip(y.view(-1), pred.view(-1)):
    confusion_matrix[t.long(), p.long()] += 1
    digit_totals[t.long()] += 1
    if t == p:
        digit_correct[t.long()] += 1

# Calcular ambas métricas
digit_accuracy = (digit_correct / digit_totals) * 100
average_accuracy = digit_accuracy.mean()
global_accuracy = (total_correct / total_samples) * 100

print("Accuracy por dígito:")
print("-" * 30)
for digit in range(10):
    print(f"Dígito {digit}: {digit_accuracy[digit]:.2f}% ({digit_totals[digit]:.0f} muestras)")
print("-" * 30)
print(f"Accuracy promedio (media de accuracies): {average_accuracy:.2f}%")
print(f"Accuracy global (total correctos/total muestras): {global_accuracy:.2f}%")
```

Epoch 1/100 Loss=0.7307 Acc=80.71%
Epoch 2/100 Loss=0.3719 Acc=89.63%
Epoch 3/100 Loss=0.3321 Acc=90.47%
Epoch 4/100 Loss=0.3143 Acc=90.88%
Epoch 5/100 Loss=0.3054 Acc=91.16%
Epoch 6/100 Loss=0.2997 Acc=91.23%
Epoch 7/100 Loss=0.2950 Acc=91.32%
Epoch 8/100 Loss=0.2923 Acc=91.53%
Epoch 9/100 Loss=0.2911 Acc=91.49%
Epoch 10/100 Loss=0.2894 Acc=91.58%
Epoch 11/100 Loss=0.2887 Acc=91.55%
Epoch 12/100 Loss=0.2874 Acc=91.56%
Epoch 13/100 Loss=0.2870 Acc=91.64%
Epoch 14/100 Loss=0.2863 Acc=91.69%
Epoch 15/100 Loss=0.2859 Acc=91.63%
Epoch 16/100 Loss=0.2845 Acc=91.66%
Epoch 17/100 Loss=0.2842 Acc=91.75%
Epoch 18/100 Loss=0.2851 Acc=91.71%
Epoch 19/100 Loss=0.2838 Acc=91.77%
Epoch 20/100 Loss=0.2850 Acc=91.65%
Epoch 21/100 Loss=0.2838 Acc=91.67%
Epoch 22/100 Loss=0.2836 Acc=91.68%
Epoch 23/100 Loss=0.2834 Acc=91.72%
Epoch 24/100 Loss=0.2833 Acc=91.73%
Epoch 25/100 Loss=0.2832 Acc=91.73%
Epoch 26/100 Loss=0.2838 Acc=91.70%
Epoch 27/100 Loss=0.2829 Acc=91.72%
Epoch 28/100 Loss=0.2833 Acc=91.70%
Epoch 29/100 Loss=0.2836 Acc=91.74%
Epoch 30/100 Loss=0.2826 Acc=91.77%
Epoch 31/100 Loss=0.2827 Acc=91.77%
Epoch 32/100 Loss=0.2834 Acc=91.71%
Epoch 33/100 Loss=0.2830 Acc=91.81%
Epoch 34/100 Loss=0.2824 Acc=91.83%
Epoch 35/100 Loss=0.2829 Acc=91.84%
Epoch 36/100 Loss=0.2832 Acc=91.73%
Epoch 37/100 Loss=0.2830 Acc=91.82%
Epoch 38/100 Loss=0.2830 Acc=91.86%
Epoch 39/100 Loss=0.2823 Acc=91.82%
Epoch 40/100 Loss=0.2821 Acc=91.83%
Epoch 41/100 Loss=0.2832 Acc=91.77%
Epoch 42/100 Loss=0.2827 Acc=91.77%
Epoch 43/100 Loss=0.2826 Acc=91.80%
Epoch 44/100 Loss=0.2830 Acc=91.83%
Epoch 45/100 Loss=0.2825 Acc=91.78%
Epoch 46/100 Loss=0.2826 Acc=91.78%
Epoch 47/100 Loss=0.2832 Acc=91.78%
Epoch 48/100 Loss=0.2813 Acc=91.83%
Epoch 49/100 Loss=0.2827 Acc=91.81%
Epoch 50/100 Loss=0.2826 Acc=91.77%
Epoch 51/100 Loss=0.2821 Acc=91.81%
Epoch 52/100 Loss=0.2823 Acc=91.74%
Epoch 53/100 Loss=0.2823 Acc=91.83%
Epoch 54/100 Loss=0.2825 Acc=91.84%
Epoch 55/100 Loss=0.2827 Acc=91.80%
Epoch 56/100 Loss=0.2821 Acc=91.84%
Epoch 57/100 Loss=0.2835 Acc=91.79%
Epoch 58/100 Loss=0.2824 Acc=91.81%
Epoch 59/100 Loss=0.2829 Acc=91.86%
Epoch 60/100 Loss=0.2821 Acc=91.83%

```
Epoch 61/100 | Loss=0.2822 | Acc=91.77%
Epoch 62/100 | Loss=0.2824 | Acc=91.85%
Epoch 63/100 | Loss=0.2825 | Acc=91.84%
Epoch 64/100 | Loss=0.2834 | Acc=91.77%
Epoch 65/100 | Loss=0.2822 | Acc=91.77%
Epoch 66/100 | Loss=0.2829 | Acc=91.71%
Epoch 67/100 | Loss=0.2834 | Acc=91.82%
Epoch 68/100 | Loss=0.2821 | Acc=91.72%
Epoch 69/100 | Loss=0.2823 | Acc=91.85%
Epoch 70/100 | Loss=0.2821 | Acc=91.78%
Epoch 71/100 | Loss=0.2820 | Acc=91.85%
Epoch 72/100 | Loss=0.2823 | Acc=91.81%
Epoch 73/100 | Loss=0.2820 | Acc=91.79%
Epoch 74/100 | Loss=0.2825 | Acc=91.75%
Epoch 75/100 | Loss=0.2826 | Acc=91.83%
Epoch 76/100 | Loss=0.2825 | Acc=91.81%
Epoch 77/100 | Loss=0.2816 | Acc=91.83%
Epoch 78/100 | Loss=0.2822 | Acc=91.79%
Epoch 79/100 | Loss=0.2830 | Acc=91.85%
Epoch 80/100 | Loss=0.2824 | Acc=91.78%
Epoch 81/100 | Loss=0.2826 | Acc=91.75%
Epoch 82/100 | Loss=0.2831 | Acc=91.76%
Epoch 83/100 | Loss=0.2821 | Acc=91.86%
Epoch 84/100 | Loss=0.2829 | Acc=91.78%
Epoch 85/100 | Loss=0.2819 | Acc=91.76%
Epoch 86/100 | Loss=0.2830 | Acc=91.84%
Epoch 87/100 | Loss=0.2812 | Acc=91.88%
Epoch 88/100 | Loss=0.2836 | Acc=91.81%
Epoch 89/100 | Loss=0.2828 | Acc=91.80%
Epoch 90/100 | Loss=0.2833 | Acc=91.80%
Epoch 91/100 | Loss=0.2825 | Acc=91.77%
Epoch 92/100 | Loss=0.2820 | Acc=91.82%
Epoch 93/100 | Loss=0.2833 | Acc=91.84%
Epoch 94/100 | Loss=0.2827 | Acc=91.78%
Epoch 95/100 | Loss=0.2823 | Acc=91.88%
Epoch 96/100 | Loss=0.2829 | Acc=91.84%
Epoch 97/100 | Loss=0.2825 | Acc=91.76%
Epoch 98/100 | Loss=0.2819 | Acc=91.88%
Epoch 99/100 | Loss=0.2826 | Acc=91.79%
Epoch 100/100 | Loss=0.2829 | Acc=91.80%
Test Accuracy: 92.09%
```

```
In [15]: class ClassifierSinEncoder(nn.Module):
    def __init__(self, num_classes=10):
        super().__init__()
        self.classifier = nn.Linear(28*28, num_classes) # conexión directa de imagen a clases

    def forward(self, x):
        x = x.view(x.size(0), -1) # aplanar la imagen
        logits = self.classifier(x)
        return logits

# Crear modelo y moverlo al dispositivo
classifier_model = ClassifierSinEncoder().to(device)
criterion_cls = nn.CrossEntropyLoss()
optimizer_cls = optim.Adam(classifier_model.parameters(), lr=1e-3)

# -----
# Entrenamiento Clasificador
# -----
```

```
num_epochs_cls = 100
for epoch in range(num_epochs_cls):
    classifier_model.train()
    running_loss = 0
    correct, total = 0, 0
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)
        optimizer_cls.zero_grad()
        outputs = classifier_model(x)
        loss = criterion_cls(outputs, y)
        loss.backward()
        optimizer_cls.step()
        running_loss += loss.item()
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
    train_acc = 100 * correct / total
    print(f"Epoch {epoch+1}/{num_epochs_cls} | Loss={running_loss/len(train_loader):.4f} | Acc={train_acc:.2f}%")

# -----
# Evaluación en Test
# -----
classifier_model.eval()
correct, total = 0, 0
with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
print(f"Test Accuracy: {100 * correct/total:.2f}%")

# Evaluación por dígito
classifier_model.eval()
confusion_matrix = torch.zeros(10, 10)
digit_totals = torch.zeros(10)
digit_correct = torch.zeros(10)
total_correct = 0
total_samples = 0

with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)

        # Contadores globales
        total_correct += (pred == y).sum().item()
        total_samples += y.size(0)

        # Contadores por dígito
        for t, p in zip(y.view(-1), pred.view(-1)):
            confusion_matrix[t.long(), p.long()] += 1
            digit_totals[t.long()] += 1
            if t == p:
                digit_correct[t.long()] += 1
```

```
# Calcular ambas métricas
digit_accuracy = (digit_correct / digit_totals) * 100
average_accuracy = digit_accuracy.mean()
global_accuracy = (total_correct / total_samples) * 100

print("Accuracy por dígito:")
print("-" * 30)
for digit in range(10):
    print(f"Dígito {digit}: {digit_accuracy[digit]:.2f}% ({digit_totals[digit]:.0f} muestras)")
print("-" * 30)
print(f"Accuracy promedio (media de accuracies): {average_accuracy:.2f}%")
print(f"Accuracy global (total correctos/total muestras): {global_accuracy:.2f}%")
```

Epoch 1/100 Loss=0.5398	Acc=87.04%
Epoch 2/100 Loss=0.3229	Acc=91.07%
Epoch 3/100 Loss=0.2953	Acc=91.79%
Epoch 4/100 Loss=0.2819	Acc=92.12%
Epoch 5/100 Loss=0.2736	Acc=92.32%
Epoch 6/100 Loss=0.2678	Acc=92.57%
Epoch 7/100 Loss=0.2634	Acc=92.61%
Epoch 8/100 Loss=0.2599	Acc=92.80%
Epoch 9/100 Loss=0.2570	Acc=92.92%
Epoch 10/100 Loss=0.2550	Acc=92.92%
Epoch 11/100 Loss=0.2523	Acc=93.01%
Epoch 12/100 Loss=0.2510	Acc=92.99%
Epoch 13/100 Loss=0.2492	Acc=93.08%
Epoch 14/100 Loss=0.2481	Acc=93.16%
Epoch 15/100 Loss=0.2464	Acc=93.19%
Epoch 16/100 Loss=0.2456	Acc=93.18%
Epoch 17/100 Loss=0.2445	Acc=93.28%
Epoch 18/100 Loss=0.2437	Acc=93.29%
Epoch 19/100 Loss=0.2422	Acc=93.31%
Epoch 20/100 Loss=0.2418	Acc=93.33%
Epoch 21/100 Loss=0.2409	Acc=93.40%
Epoch 22/100 Loss=0.2403	Acc=93.37%
Epoch 23/100 Loss=0.2396	Acc=93.42%
Epoch 24/100 Loss=0.2387	Acc=93.45%
Epoch 25/100 Loss=0.2382	Acc=93.45%
Epoch 26/100 Loss=0.2378	Acc=93.45%
Epoch 27/100 Loss=0.2375	Acc=93.45%
Epoch 28/100 Loss=0.2365	Acc=93.52%
Epoch 29/100 Loss=0.2364	Acc=93.52%
Epoch 30/100 Loss=0.2357	Acc=93.48%
Epoch 31/100 Loss=0.2353	Acc=93.51%
Epoch 32/100 Loss=0.2347	Acc=93.57%
Epoch 33/100 Loss=0.2342	Acc=93.55%
Epoch 34/100 Loss=0.2337	Acc=93.55%
Epoch 35/100 Loss=0.2336	Acc=93.56%
Epoch 36/100 Loss=0.2329	Acc=93.58%
Epoch 37/100 Loss=0.2326	Acc=93.53%
Epoch 38/100 Loss=0.2327	Acc=93.63%
Epoch 39/100 Loss=0.2318	Acc=93.58%
Epoch 40/100 Loss=0.2317	Acc=93.58%
Epoch 41/100 Loss=0.2314	Acc=93.63%
Epoch 42/100 Loss=0.2315	Acc=93.64%
Epoch 43/100 Loss=0.2311	Acc=93.59%
Epoch 44/100 Loss=0.2303	Acc=93.69%
Epoch 45/100 Loss=0.2302	Acc=93.63%
Epoch 46/100 Loss=0.2299	Acc=93.67%
Epoch 47/100 Loss=0.2296	Acc=93.67%
Epoch 48/100 Loss=0.2297	Acc=93.65%
Epoch 49/100 Loss=0.2288	Acc=93.66%
Epoch 50/100 Loss=0.2292	Acc=93.69%
Epoch 51/100 Loss=0.2287	Acc=93.67%
Epoch 52/100 Loss=0.2287	Acc=93.69%
Epoch 53/100 Loss=0.2284	Acc=93.71%
Epoch 54/100 Loss=0.2281	Acc=93.70%
Epoch 55/100 Loss=0.2278	Acc=93.68%
Epoch 56/100 Loss=0.2277	Acc=93.70%
Epoch 57/100 Loss=0.2275	Acc=93.67%
Epoch 58/100 Loss=0.2273	Acc=93.66%
Epoch 59/100 Loss=0.2273	Acc=93.76%
Epoch 60/100 Loss=0.2267	Acc=93.69%

```
Epoch 61/100 | Loss=0.2268 | Acc=93.75%
Epoch 62/100 | Loss=0.2262 | Acc=93.75%
Epoch 63/100 | Loss=0.2267 | Acc=93.74%
Epoch 64/100 | Loss=0.2263 | Acc=93.77%
Epoch 65/100 | Loss=0.2261 | Acc=93.73%
Epoch 66/100 | Loss=0.2257 | Acc=93.74%
Epoch 67/100 | Loss=0.2257 | Acc=93.70%
Epoch 68/100 | Loss=0.2258 | Acc=93.69%
Epoch 69/100 | Loss=0.2252 | Acc=93.79%
Epoch 70/100 | Loss=0.2251 | Acc=93.73%
Epoch 71/100 | Loss=0.2247 | Acc=93.76%
Epoch 72/100 | Loss=0.2247 | Acc=93.81%
Epoch 73/100 | Loss=0.2252 | Acc=93.80%
Epoch 74/100 | Loss=0.2245 | Acc=93.78%
Epoch 75/100 | Loss=0.2246 | Acc=93.81%
Epoch 76/100 | Loss=0.2240 | Acc=93.85%
Epoch 77/100 | Loss=0.2243 | Acc=93.77%
Epoch 78/100 | Loss=0.2238 | Acc=93.80%
Epoch 79/100 | Loss=0.2242 | Acc=93.77%
Epoch 80/100 | Loss=0.2235 | Acc=93.80%
Epoch 81/100 | Loss=0.2235 | Acc=93.83%
Epoch 82/100 | Loss=0.2235 | Acc=93.81%
Epoch 83/100 | Loss=0.2238 | Acc=93.90%
Epoch 84/100 | Loss=0.2237 | Acc=93.86%
Epoch 85/100 | Loss=0.2228 | Acc=93.87%
Epoch 86/100 | Loss=0.2229 | Acc=93.85%
Epoch 87/100 | Loss=0.2230 | Acc=93.81%
Epoch 88/100 | Loss=0.2228 | Acc=93.80%
Epoch 89/100 | Loss=0.2227 | Acc=93.84%
Epoch 90/100 | Loss=0.2225 | Acc=93.84%
Epoch 91/100 | Loss=0.2230 | Acc=93.83%
Epoch 92/100 | Loss=0.2226 | Acc=93.80%
Epoch 93/100 | Loss=0.2218 | Acc=93.83%
Epoch 94/100 | Loss=0.2223 | Acc=93.81%
Epoch 95/100 | Loss=0.2219 | Acc=93.92%
Epoch 96/100 | Loss=0.2219 | Acc=93.84%
Epoch 97/100 | Loss=0.2218 | Acc=93.85%
Epoch 98/100 | Loss=0.2223 | Acc=93.79%
Epoch 99/100 | Loss=0.2216 | Acc=93.86%
Epoch 100/100 | Loss=0.2213 | Acc=93.86%
```

Test Accuracy: 92.73%

Accuracy por dígito:

```
-----  
Dígito 0: 96.94% (980 muestras)  
Dígito 1: 97.53% (1135 muestras)  
Dígito 2: 89.73% (1032 muestras)  
Dígito 3: 90.10% (1010 muestras)  
Dígito 4: 92.77% (982 muestras)  
Dígito 5: 87.44% (892 muestras)  
Dígito 6: 95.93% (958 muestras)  
Dígito 7: 92.41% (1028 muestras)  
Dígito 8: 90.97% (974 muestras)  
Dígito 9: 92.57% (1009 muestras)
```

Accuracy promedio (media de accuracies): 92.64%

Accuracy global (total correctos/total muestras): 92.73%

```
In [16]: class ClassifierFromEncoder(nn.Module):
    def __init__(self, autoencoder, num_classes=10):
        super().__init__()
```

```
# Capas del encoder congeladas
self.encoder_fc1 = autoencoder.fc1
self.encoder_fc2 = autoencoder.fc2
for param in self.encoder_fc1.parameters():
    param.requires_grad = False
for param in self.encoder_fc2.parameters():
    param.requires_grad = False

# Capas del clasificador
self.classifier1 = nn.Linear(156, 156) # Nueva capa oculta
self.classifier2 = nn.Linear(156, num_classes) # Capa de salida

def forward(self, x):
    x = x.view(x.size(0), -1)
    # Encoder congelado
    with torch.no_grad():
        x = F.relu(self.encoder_fc1(x))
        encoded = F.relu(self.encoder_fc2(x))
    # Clasificador entrenable
    x = F.relu(self.classifier1(encoded))
    logits = self.classifier2(x)
    return logits

classifier_model = ClassifierFromEncoder(model_cnn).to(device)
criterion_cls = nn.CrossEntropyLoss()
# Optimizador solo con parámetros de las capas classifier
optimizer_cls = optim.Adam([
    *classifier_model.classifier1.parameters(),
    *classifier_model.classifier2.parameters()
], lr=1e-3)
# -----
# Entrenamiento Clasificador
# -----
num_epochs_cls = 100
for epoch in range(num_epochs_cls):
    classifier_model.train()
    running_loss = 0
    correct, total = 0, 0
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)
        optimizer_cls.zero_grad()
        outputs = classifier_model(x)
        loss = criterion_cls(outputs, y)
        loss.backward()
        optimizer_cls.step()
        running_loss += loss.item()
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
    train_acc = 100 * correct / total
    print(f"Epoch {epoch+1}/{num_epochs_cls} | Loss={running_loss/len(train_loader):.4f} | Acc={train_acc:.2f}%")

# -----
# Evaluación en Test
# -----
classifier_model.eval()
correct, total = 0, 0
with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
```

```
outputs = classifier_model(x)
_, pred = outputs.max(1)
correct += pred.eq(y).sum().item()
total += y.size(0)
print(f"Test Accuracy: {100 * correct/total:.2f}%")

# Evaluación por dígito
classifier_model.eval()
confusion_matrix = torch.zeros(10, 10)
digit_totals = torch.zeros(10)
digit_correct = torch.zeros(10)
total_correct = 0
total_samples = 0

with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)

        # Contadores globales
        total_correct += (pred == y).sum().item()
        total_samples += y.size(0)

        # Contadores por dígito
        for t, p in zip(y.view(-1), pred.view(-1)):
            confusion_matrix[t.long(), p.long()] += 1
            digit_totals[t.long()] += 1
            if t == p:
                digit_correct[t.long()] += 1

# Calcular ambas métricas
digit_accuracy = (digit_correct / digit_totals) * 100
average_accuracy = digit_accuracy.mean()
global_accuracy = (total_correct / total_samples) * 100

print("Accuracy por dígito:")
print("-" * 30)
for digit in range(10):
    print(f"Dígito {digit}: {digit_accuracy[digit]:.2f}% ({digit_totals[digit]:.0f} muestras)")
print("-" * 30)
print(f"Accuracy promedio (media de accuracies): {average_accuracy:.2f}%")
print(f"Accuracy global (total correctos/total muestras): {global_accuracy:.2f}%")
```

Epoch 1/100 Loss=0.4646	Acc=86.90%
Epoch 2/100 Loss=0.2832	Acc=91.65%
Epoch 3/100 Loss=0.2215	Acc=93.48%
Epoch 4/100 Loss=0.1733	Acc=94.85%
Epoch 5/100 Loss=0.1412	Acc=95.86%
Epoch 6/100 Loss=0.1200	Acc=96.42%
Epoch 7/100 Loss=0.1059	Acc=96.73%
Epoch 8/100 Loss=0.0955	Acc=97.11%
Epoch 9/100 Loss=0.0846	Acc=97.47%
Epoch 10/100 Loss=0.0781	Acc=97.62%
Epoch 11/100 Loss=0.0730	Acc=97.72%
Epoch 12/100 Loss=0.0660	Acc=97.93%
Epoch 13/100 Loss=0.0634	Acc=98.02%
Epoch 14/100 Loss=0.0583	Acc=98.12%
Epoch 15/100 Loss=0.0547	Acc=98.26%
Epoch 16/100 Loss=0.0519	Acc=98.33%
Epoch 17/100 Loss=0.0479	Acc=98.50%
Epoch 18/100 Loss=0.0467	Acc=98.52%
Epoch 19/100 Loss=0.0432	Acc=98.60%
Epoch 20/100 Loss=0.0412	Acc=98.59%
Epoch 21/100 Loss=0.0390	Acc=98.71%
Epoch 22/100 Loss=0.0366	Acc=98.80%
Epoch 23/100 Loss=0.0379	Acc=98.74%
Epoch 24/100 Loss=0.0334	Acc=98.92%
Epoch 25/100 Loss=0.0335	Acc=98.88%
Epoch 26/100 Loss=0.0305	Acc=99.00%
Epoch 27/100 Loss=0.0295	Acc=99.01%
Epoch 28/100 Loss=0.0271	Acc=99.08%
Epoch 29/100 Loss=0.0265	Acc=99.11%
Epoch 30/100 Loss=0.0248	Acc=99.17%
Epoch 31/100 Loss=0.0247	Acc=99.19%
Epoch 32/100 Loss=0.0235	Acc=99.26%
Epoch 33/100 Loss=0.0224	Acc=99.24%
Epoch 34/100 Loss=0.0231	Acc=99.22%
Epoch 35/100 Loss=0.0202	Acc=99.30%
Epoch 36/100 Loss=0.0202	Acc=99.30%
Epoch 37/100 Loss=0.0165	Acc=99.45%
Epoch 38/100 Loss=0.0174	Acc=99.38%
Epoch 39/100 Loss=0.0195	Acc=99.32%
Epoch 40/100 Loss=0.0160	Acc=99.46%
Epoch 41/100 Loss=0.0146	Acc=99.51%
Epoch 42/100 Loss=0.0171	Acc=99.44%
Epoch 43/100 Loss=0.0168	Acc=99.42%
Epoch 44/100 Loss=0.0139	Acc=99.52%
Epoch 45/100 Loss=0.0164	Acc=99.41%
Epoch 46/100 Loss=0.0150	Acc=99.52%
Epoch 47/100 Loss=0.0127	Acc=99.61%
Epoch 48/100 Loss=0.0143	Acc=99.52%
Epoch 49/100 Loss=0.0138	Acc=99.55%
Epoch 50/100 Loss=0.0115	Acc=99.62%
Epoch 51/100 Loss=0.0134	Acc=99.52%
Epoch 52/100 Loss=0.0107	Acc=99.64%
Epoch 53/100 Loss=0.0125	Acc=99.56%
Epoch 54/100 Loss=0.0117	Acc=99.61%
Epoch 55/100 Loss=0.0139	Acc=99.51%
Epoch 56/100 Loss=0.0075	Acc=99.77%
Epoch 57/100 Loss=0.0127	Acc=99.56%
Epoch 58/100 Loss=0.0115	Acc=99.60%
Epoch 59/100 Loss=0.0109	Acc=99.63%
Epoch 60/100 Loss=0.0091	Acc=99.67%

```
Epoch 61/100 | Loss=0.0141 | Acc=99.48%
Epoch 62/100 | Loss=0.0079 | Acc=99.72%
Epoch 63/100 | Loss=0.0070 | Acc=99.75%
Epoch 64/100 | Loss=0.0134 | Acc=99.52%
Epoch 65/100 | Loss=0.0083 | Acc=99.71%
Epoch 66/100 | Loss=0.0090 | Acc=99.69%
Epoch 67/100 | Loss=0.0111 | Acc=99.63%
Epoch 68/100 | Loss=0.0107 | Acc=99.59%
Epoch 69/100 | Loss=0.0075 | Acc=99.74%
Epoch 70/100 | Loss=0.0105 | Acc=99.62%
Epoch 71/100 | Loss=0.0073 | Acc=99.74%
Epoch 72/100 | Loss=0.0124 | Acc=99.57%
Epoch 73/100 | Loss=0.0088 | Acc=99.68%
Epoch 74/100 | Loss=0.0068 | Acc=99.76%
Epoch 75/100 | Loss=0.0079 | Acc=99.70%
Epoch 76/100 | Loss=0.0091 | Acc=99.70%
Epoch 77/100 | Loss=0.0096 | Acc=99.67%
Epoch 78/100 | Loss=0.0058 | Acc=99.79%
Epoch 79/100 | Loss=0.0081 | Acc=99.72%
Epoch 80/100 | Loss=0.0089 | Acc=99.69%
Epoch 81/100 | Loss=0.0094 | Acc=99.67%
Epoch 82/100 | Loss=0.0079 | Acc=99.72%
Epoch 83/100 | Loss=0.0067 | Acc=99.76%
Epoch 84/100 | Loss=0.0089 | Acc=99.69%
Epoch 85/100 | Loss=0.0066 | Acc=99.76%
Epoch 86/100 | Loss=0.0073 | Acc=99.75%
Epoch 87/100 | Loss=0.0050 | Acc=99.84%
Epoch 88/100 | Loss=0.0091 | Acc=99.66%
Epoch 89/100 | Loss=0.0093 | Acc=99.67%
Epoch 90/100 | Loss=0.0032 | Acc=99.90%
Epoch 91/100 | Loss=0.0135 | Acc=99.57%
Epoch 92/100 | Loss=0.0066 | Acc=99.78%
Epoch 93/100 | Loss=0.0064 | Acc=99.78%
Epoch 94/100 | Loss=0.0070 | Acc=99.77%
Epoch 95/100 | Loss=0.0063 | Acc=99.77%
Epoch 96/100 | Loss=0.0084 | Acc=99.69%
Epoch 97/100 | Loss=0.0013 | Acc=99.97%
Epoch 98/100 | Loss=0.0130 | Acc=99.58%
Epoch 99/100 | Loss=0.0070 | Acc=99.77%
Epoch 100/100 | Loss=0.0067 | Acc=99.77%
Test Accuracy: 97.94%
```

Accuracy por dígito:

```
-----
Dígito 0: 99.08% (980 muestras)
Dígito 1: 99.65% (1135 muestras)
Dígito 2: 97.87% (1032 muestras)
Dígito 3: 98.71% (1010 muestras)
Dígito 4: 98.27% (982 muestras)
Dígito 5: 96.19% (892 muestras)
Dígito 6: 98.23% (958 muestras)
Dígito 7: 97.28% (1028 muestras)
Dígito 8: 97.64% (974 muestras)
Dígito 9: 96.13% (1009 muestras)
```

```
Accuracy promedio (media de accuracies): 97.90%
Accuracy global (total correctos/total muestras): 97.94%
```

```
In [17]: class ClassifierFromEncoder(nn.Module):
    def __init__(self, autoencoder, num_classes=10):
        super().__init__()
```

```
# Solo copiamos las capas, sin congelarlas
self.encoder_fc1 = autoencoder.fc1
self.encoder_fc2 = autoencoder.fc2
self.classifier = nn.Linear(156, num_classes)

def forward(self, x):
    x = x.view(x.size(0), -1)
    # Sin torch.no_grad()
    x = F.relu(self.encoder_fc1(x))
    encoded = F.relu(self.encoder_fc2(x))
    logits = self.classifier(encoded)
    return logits

classifier_model = ClassifierFromEncoder(model_cnn).to(device)
criterion_cls = nn.CrossEntropyLoss()
# Optimizador con todos los parámetros
optimizer_cls = optim.Adam(classifier_model.parameters(), lr=1e-3)

# -----
# Entrenamiento Clasificador
# -----
num_epochs_cls = 100
for epoch in range(num_epochs_cls):
    classifier_model.train()
    running_loss = 0
    correct, total = 0, 0
    for x, y in train_loader:
        x, y = x.to(device), y.to(device)
        optimizer_cls.zero_grad()
        outputs = classifier_model(x)
        loss = criterion_cls(outputs, y)
        loss.backward()
        optimizer_cls.step()
        running_loss += loss.item()
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
    train_acc = 100 * correct / total
    print(f"Epoch {epoch+1}/{num_epochs_cls} | Loss={running_loss/len(train_loader):.4f} | Acc={train_acc:.2f}%")

# -----
# Evaluación en Test
# -----
classifier_model.eval()
correct, total = 0, 0
with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)
        correct += pred.eq(y).sum().item()
        total += y.size(0)
print(f"Test Accuracy: {100 * correct/total:.2f}%")

# Evaluación por dígito
classifier_model.eval()
confusion_matrix = torch.zeros(10, 10)
```

```
digit_totals = torch.zeros(10)
digit_correct = torch.zeros(10)
total_correct = 0
total_samples = 0

with torch.no_grad():
    for x, y in test_loader:
        x, y = x.to(device), y.to(device)
        outputs = classifier_model(x)
        _, pred = outputs.max(1)

        # Contadores globales
        total_correct += (pred == y).sum().item()
        total_samples += y.size(0)

        # Contadores por dígito
        for t, p in zip(y.view(-1), pred.view(-1)):
            confusion_matrix[t.long(), p.long()] += 1
            digit_totals[t.long()] += 1
            if t == p:
                digit_correct[t.long()] += 1

# Calcular ambas métricas
digit_accuracy = (digit_correct / digit_totals) * 100
average_accuracy = digit_accuracy.mean()
global_accuracy = (total_correct / total_samples) * 100

print("Accuracy por dígito:")
print("-" * 30)
for digit in range(10):
    print(f"Dígito {digit}: {digit_accuracy[digit]:.2f}% ({digit_totals[digit]:.0f} muestras)")
print("-" * 30)
print(f"Accuracy promedio (media de accuracies): {average_accuracy:.2f}%")
print(f"Accuracy global (total correctos/total muestras): {global_accuracy:.2f}%")
```

Epoch 1/100 Loss=0.7107 Acc=81.59%
Epoch 2/100 Loss=0.3689 Acc=89.73%
Epoch 3/100 Loss=0.3295 Acc=90.53%
Epoch 4/100 Loss=0.3135 Acc=90.94%
Epoch 5/100 Loss=0.3042 Acc=91.09%
Epoch 6/100 Loss=0.2988 Acc=91.26%
Epoch 7/100 Loss=0.2950 Acc=91.42%
Epoch 8/100 Loss=0.2922 Acc=91.55%
Epoch 9/100 Loss=0.2899 Acc=91.47%
Epoch 10/100 Loss=0.2894 Acc=91.53%
Epoch 11/100 Loss=0.2879 Acc=91.56%
Epoch 12/100 Loss=0.2875 Acc=91.63%
Epoch 13/100 Loss=0.2855 Acc=91.65%
Epoch 14/100 Loss=0.2854 Acc=91.61%
Epoch 15/100 Loss=0.2857 Acc=91.71%
Epoch 16/100 Loss=0.2855 Acc=91.64%
Epoch 17/100 Loss=0.2849 Acc=91.71%
Epoch 18/100 Loss=0.2841 Acc=91.69%
Epoch 19/100 Loss=0.2833 Acc=91.80%
Epoch 20/100 Loss=0.2836 Acc=91.77%
Epoch 21/100 Loss=0.2840 Acc=91.77%
Epoch 22/100 Loss=0.2831 Acc=91.78%
Epoch 23/100 Loss=0.2838 Acc=91.70%
Epoch 24/100 Loss=0.2834 Acc=91.75%
Epoch 25/100 Loss=0.2836 Acc=91.71%
Epoch 26/100 Loss=0.2834 Acc=91.70%
Epoch 27/100 Loss=0.2831 Acc=91.82%
Epoch 28/100 Loss=0.2824 Acc=91.81%
Epoch 29/100 Loss=0.2828 Acc=91.69%
Epoch 30/100 Loss=0.2824 Acc=91.76%
Epoch 31/100 Loss=0.2828 Acc=91.74%
Epoch 32/100 Loss=0.2828 Acc=91.74%
Epoch 33/100 Loss=0.2840 Acc=91.75%
Epoch 34/100 Loss=0.2836 Acc=91.81%
Epoch 35/100 Loss=0.2824 Acc=91.83%
Epoch 36/100 Loss=0.2825 Acc=91.80%
Epoch 37/100 Loss=0.2828 Acc=91.82%
Epoch 38/100 Loss=0.2840 Acc=91.78%
Epoch 39/100 Loss=0.2819 Acc=91.83%
Epoch 40/100 Loss=0.2819 Acc=91.80%
Epoch 41/100 Loss=0.2832 Acc=91.78%
Epoch 42/100 Loss=0.2819 Acc=91.84%
Epoch 43/100 Loss=0.2826 Acc=91.78%
Epoch 44/100 Loss=0.2833 Acc=91.74%
Epoch 45/100 Loss=0.2826 Acc=91.72%
Epoch 46/100 Loss=0.2824 Acc=91.88%
Epoch 47/100 Loss=0.2816 Acc=91.84%
Epoch 48/100 Loss=0.2827 Acc=91.83%
Epoch 49/100 Loss=0.2831 Acc=91.75%
Epoch 50/100 Loss=0.2818 Acc=91.83%
Epoch 51/100 Loss=0.2824 Acc=91.81%
Epoch 52/100 Loss=0.2831 Acc=91.75%
Epoch 53/100 Loss=0.2826 Acc=91.81%
Epoch 54/100 Loss=0.2822 Acc=91.73%
Epoch 55/100 Loss=0.2825 Acc=91.78%
Epoch 56/100 Loss=0.2821 Acc=91.79%
Epoch 57/100 Loss=0.2835 Acc=91.74%
Epoch 58/100 Loss=0.2826 Acc=91.79%
Epoch 59/100 Loss=0.2826 Acc=91.81%
Epoch 60/100 Loss=0.2824 Acc=91.69%

```
Epoch 61/100 | Loss=0.2827 | Acc=91.82%
Epoch 62/100 | Loss=0.2829 | Acc=91.80%
Epoch 63/100 | Loss=0.2829 | Acc=91.84%
Epoch 64/100 | Loss=0.2825 | Acc=91.80%
Epoch 65/100 | Loss=0.2824 | Acc=91.86%
Epoch 66/100 | Loss=0.2831 | Acc=91.70%
Epoch 67/100 | Loss=0.2825 | Acc=91.70%
Epoch 68/100 | Loss=0.2822 | Acc=91.85%
Epoch 69/100 | Loss=0.2824 | Acc=91.84%
Epoch 70/100 | Loss=0.2827 | Acc=91.87%
Epoch 71/100 | Loss=0.2826 | Acc=91.83%
Epoch 72/100 | Loss=0.2816 | Acc=91.89%
Epoch 73/100 | Loss=0.2822 | Acc=91.77%
Epoch 74/100 | Loss=0.2826 | Acc=91.81%
Epoch 75/100 | Loss=0.2833 | Acc=91.82%
Epoch 76/100 | Loss=0.2829 | Acc=91.78%
Epoch 77/100 | Loss=0.2817 | Acc=91.81%
Epoch 78/100 | Loss=0.2819 | Acc=91.85%
Epoch 79/100 | Loss=0.2825 | Acc=91.75%
Epoch 80/100 | Loss=0.2820 | Acc=91.83%
Epoch 81/100 | Loss=0.2817 | Acc=91.80%
Epoch 82/100 | Loss=0.2818 | Acc=91.81%
Epoch 83/100 | Loss=0.2826 | Acc=91.84%
Epoch 84/100 | Loss=0.2828 | Acc=91.78%
Epoch 85/100 | Loss=0.2820 | Acc=91.81%
Epoch 86/100 | Loss=0.2820 | Acc=91.81%
Epoch 87/100 | Loss=0.2827 | Acc=91.82%
Epoch 88/100 | Loss=0.2822 | Acc=91.80%
Epoch 89/100 | Loss=0.2819 | Acc=91.82%
Epoch 90/100 | Loss=0.2825 | Acc=91.81%
Epoch 91/100 | Loss=0.2831 | Acc=91.80%
Epoch 92/100 | Loss=0.2826 | Acc=91.82%
Epoch 93/100 | Loss=0.2831 | Acc=91.79%
Epoch 94/100 | Loss=0.2819 | Acc=91.79%
Epoch 95/100 | Loss=0.2826 | Acc=91.73%
Epoch 96/100 | Loss=0.2819 | Acc=91.84%
Epoch 97/100 | Loss=0.2820 | Acc=91.80%
Epoch 98/100 | Loss=0.2820 | Acc=91.82%
Epoch 99/100 | Loss=0.2830 | Acc=91.75%
Epoch 100/100 | Loss=0.2818 | Acc=91.80%
```

Test Accuracy: 92.35%

Accuracy por dígito:

```
-----
Dígito 0: 97.65% (980 muestras)
Dígito 1: 97.97% (1135 muestras)
Dígito 2: 90.41% (1032 muestras)
Dígito 3: 90.89% (1010 muestras)
Dígito 4: 93.28% (982 muestras)
Dígito 5: 85.87% (892 muestras)
Dígito 6: 95.09% (958 muestras)
Dígito 7: 91.54% (1028 muestras)
Dígito 8: 89.12% (974 muestras)
Dígito 9: 90.49% (1009 muestras)
```

Accuracy promedio (media de accuracies): 92.23%

Accuracy global (total correctos/total muestras): 92.35%